

Relational Algebra and Query Languages

Venkatesh Vinayakarao

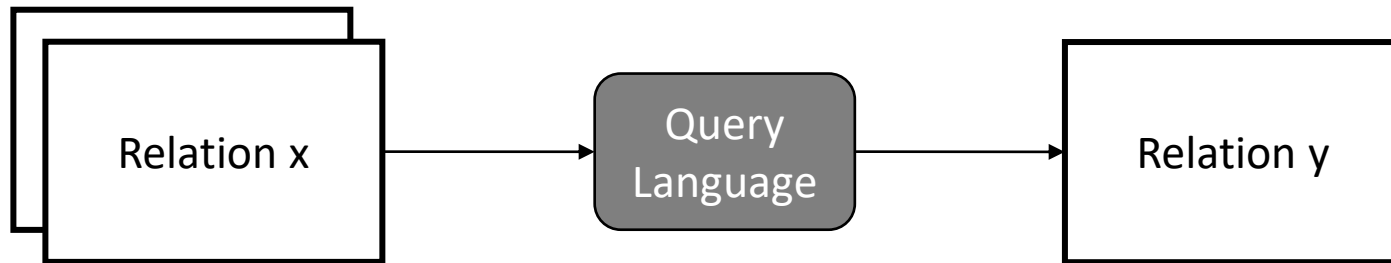
venkateshv@cmi.ac.in

<http://vvtesh.co.in>

Chennai Mathematical Institute

Relational Algebra and Query Languages

Query Languages



Procedural Language

Relational Algebra

Declarative Languages

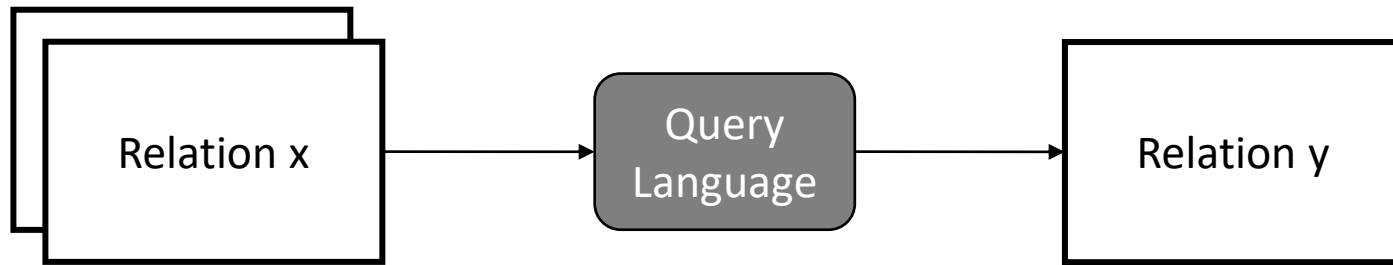
Tuple Relational Calculus

Domain Relational Calculus

Popular Language

SQL

Relational Algebra



Fundamental Operations

select, project, rename
set difference, cartesian product, union

More Operations

set intersection, natural join, assignment

Select Operation

- Notation: $\sigma_p(r)$ where p is called the selection predicate

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

(a) The *instructor* table

- Example of selection:

$\sigma_{\text{dept_name}=\text{"Physics"}}(\text{instructor})$

Select Operation

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

(a) The *instructor* table

- Example of selection:

$\sigma_{\text{dept_name}=\text{"Physics"}}(\text{instructor})$

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 33456 | Gold | Physics | 87000 |

Project Operation

- Notation: $\Pi_{A_1, A_2, \dots, A_k}(r)$ where A_i are attribute names

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

(a) The *instructor* table

- Example of projection:

$$\Pi_{ID, name, salary}(instructor)$$

- Duplicate rows removed from result, since relations are sets

Projection

$\Pi_{ID, name, salary}(instructor)$

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

(a) The *instructor* table



| <i>ID</i> | <i>name</i> | <i>salary</i> |
|-----------|-------------|---------------|
| 10101 | Srinivasan | 65000 |
| 12121 | Wu | 90000 |
| 15151 | Mozart | 40000 |
| 22222 | Einstein | 95000 |
| 32343 | El Said | 60000 |
| 33456 | Gold | 87000 |
| 45565 | Katz | 75000 |
| 58583 | Califieri | 62000 |
| 76543 | Singh | 80000 |
| 76766 | Crick | 72000 |
| 83821 | Brandt | 92000 |
| 98345 | Kim | 80000 |

Union Operation

- Notation: $r \cup s$. Defined as:

$$r \cup s = \{t \mid t \in r \text{ or } t \in s\}$$

- For $r \cup s$ to be valid:
 - r, s must have the *same arity* (same number of attributes)
 - The attribute domains must be *compatible* (example: 2nd column of r deals with the same type of values as does the 2nd column of s)
- Example: to find all courses taught in the Fall 2009 semester, or in the Spring 2010 semester, or in both

$\Pi_{course_id} (\sigma_{semester="Fall" \wedge year=2009} (section)) \cup$

$\Pi_{course_id} (\sigma_{semester="Spring" \wedge year=2010} (section))$

Union, Selection and Projection

- $\Pi_{course_id} (\sigma_{semester="Fall" \wedge year=2009}(section)) \cup$
 $\Pi_{course_id} (\sigma_{semester="Spring" \wedge year=2010}(section))$

| course_id | sec_id | semester | year | building | room_number | time_slot_id |
|-----------|--------|----------|------|----------|-------------|--------------|
| BIO-101 | 1 | Summer | 2009 | Painter | 514 | B |
| BIO-301 | 1 | Summer | 2010 | Painter | 514 | A |
| CS-101 | 1 | Fall | 2009 | Packard | 101 | H |
| CS-101 | 1 | Spring | 2010 | Packard | 101 | F |
| CS-190 | 1 | Spring | 2009 | Taylor | 3128 | E |
| CS-190 | 2 | Spring | 2009 | Taylor | 3128 | A |
| CS-315 | 1 | Spring | 2010 | Watson | 120 | D |
| CS-319 | 1 | Spring | 2010 | Watson | 100 | B |
| CS-319 | 2 | Spring | 2010 | Taylor | 3128 | C |
| CS-347 | 1 | Fall | 2009 | Taylor | 3128 | A |
| EE-181 | 1 | Spring | 2009 | Taylor | 3128 | C |
| FIN-201 | 1 | Spring | 2010 | Packard | 101 | B |
| HIS-351 | 1 | Spring | 2010 | Painter | 514 | C |
| MU-199 | 1 | Spring | 2010 | Packard | 101 | D |
| PHY-101 | 1 | Fall | 2009 | Watson | 100 | A |



| course_id |
|-----------|
| CS-101 |
| CS-315 |
| CS-319 |
| CS-347 |
| FIN-201 |
| HIS-351 |
| MU-199 |
| PHY-101 |

Set Difference Operation

- Notation: $r - s$. Defined as

$$r - s = \{t \mid t \in r \textbf{ and } t \notin s\}$$

- Example: to find all courses taught in the Fall 2009 semester, but not in the Spring 2010 semester

$$\Pi_{course_id} (\sigma_{semester="Fall" \wedge year=2009}(section)) - \Pi_{course_id} (\sigma_{semester="Spring" \wedge year=2010}(section))$$

Union, Selection and Projection

- $\Pi_{course_id} (\sigma_{semester="Fall" \wedge year=2009}(section)) - \Pi_{course_id} (\sigma_{semester="Spring" \wedge year=2010}(section))$

| course_id | sec_id | semester | year | building | room_number | time_slot_id |
|-----------|--------|----------|------|----------|-------------|--------------|
| BIO-101 | 1 | Summer | 2009 | Painter | 514 | B |
| BIO-301 | 1 | Summer | 2010 | Painter | 514 | A |
| CS-101 | 1 | Fall | 2009 | Packard | 101 | H |
| CS-101 | 1 | Spring | 2010 | Packard | 101 | F |
| CS-190 | 1 | Spring | 2009 | Taylor | 3128 | E |
| CS-190 | 2 | Spring | 2009 | Taylor | 3128 | A |
| CS-315 | 1 | Spring | 2010 | Watson | 120 | D |
| CS-319 | 1 | Spring | 2010 | Watson | 100 | B |
| CS-319 | 2 | Spring | 2010 | Taylor | 3128 | C |
| CS-347 | 1 | Fall | 2009 | Taylor | 3128 | A |
| EE-181 | 1 | Spring | 2009 | Taylor | 3128 | C |
| FIN-201 | 1 | Spring | 2010 | Packard | 101 | B |
| HIS-351 | 1 | Spring | 2010 | Painter | 514 | C |
| MU-199 | 1 | Spring | 2010 | Packard | 101 | D |
| PHY-101 | 1 | Fall | 2009 | Watson | 100 | A |



| course_id |
|-----------|
| CS-347 |
| PHY-101 |

Set-Intersection Operation

- Notation: $r \cap s$
- Defined as:
- $r \cap s = \{ t \mid t \in r \textbf{ and } t \in s \}$

Quiz: True/False?

$$r \cap s = r - (r - s)$$

Cartesian-Product Operation

- Notation $r \times s$


| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 12121 | Wu | Finance | 90000 |
| 15151 | Mozart | Music | 40000 |
| 22222 | Einstein | Physics | 95000 |
| 32343 | El Said | History | 60000 |
| 33456 | Gold | Physics | 87000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 58583 | Califieri | History | 62000 |
| 76543 | Singh | Finance | 80000 |
| 76766 | Crick | Biology | 72000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 98345 | Kim | Elec. Eng. | 80000 |

instructor relation

| <i>ID</i> | <i>course_id</i> | <i>sec_id</i> | <i>semester</i> | <i>year</i> |
|-----------|------------------|---------------|-----------------|-------------|
| 10101 | CS-101 | 1 | Fall | 2009 |
| 10101 | CS-315 | 1 | Spring | 2010 |
| 10101 | CS-347 | 1 | Fall | 2009 |
| 12121 | FIN-201 | 1 | Spring | 2010 |
| 15151 | MU-199 | 1 | Spring | 2010 |
| 22222 | PHY-101 | 1 | Fall | 2009 |
| 32343 | HIS-351 | 1 | Spring | 2010 |
| 45565 | CS-101 | 1 | Spring | 2010 |
| 45565 | CS-319 | 1 | Spring | 2010 |
| 76766 | BIO-101 | 1 | Summer | 2009 |
| 76766 | BIO-301 | 1 | Summer | 2010 |
| 83821 | CS-190 | 1 | Spring | 2009 |
| 83821 | CS-190 | 2 | Spring | 2009 |
| 83821 | CS-319 | 2 | Spring | 2010 |
| 98345 | EE-181 | 1 | Spring | 2009 |

teaches relation

The Fly on the Ceiling - Descartes



French mathematician René
Descartes (1596-1650)

Image Source:

<http://sites.psu.edu/solvingproblems>

<https://wild.maths.org/ren%C3%A9-descartes-and-fly-ceiling>

vectorstock.com

The *instructor* X teaches table

[illegible]

Join Operation

- The Cartesian-Product *instructor X teaches* associates every tuple of instructor with every tuple of teaches.
- Most of the resulting rows have information about instructors who did NOT teach a particular course.
- To get only those tuples of *instructor X teaches* that pertain to instructors and the courses that they taught, we write:

$$\sigma_{instructor.id = teaches.id} (instructor \times teaches)$$

Join Operation (Cont.)

- $\sigma_{instructor.id = teaches.id} (instructor \times teaches)$

| <i>instructor.ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> | <i>teaches.ID</i> | <i>course_id</i> | <i>sec_id</i> | <i>semester</i> | <i>year</i> |
|----------------------|-------------|------------------|---------------|-------------------|------------------|---------------|-----------------|-------------|
| 10101 | Srinivasan | Comp. Sci. | 65000 | 10101 | CS-101 | 1 | Fall | 2017 |
| 10101 | Srinivasan | Comp. Sci. | 65000 | 10101 | CS-315 | 1 | Spring | 2018 |
| 10101 | Srinivasan | Comp. Sci. | 65000 | 10101 | CS-347 | 1 | Fall | 2017 |
| 12121 | Wu | Finance | 90000 | 12121 | FIN-201 | 1 | Spring | 2018 |
| 15151 | Mozart | Music | 40000 | 15151 | MU-199 | 1 | Spring | 2018 |
| 22222 | Einstein | Physics | 95000 | 22222 | PHY-101 | 1 | Fall | 2017 |
| 32343 | El Said | History | 60000 | 32343 | HIS-351 | 1 | Spring | 2018 |
| 45565 | Katz | Comp. Sci. | 75000 | 45565 | CS-101 | 1 | Spring | 2018 |
| 45565 | Katz | Comp. Sci. | 75000 | 45565 | CS-319 | 1 | Spring | 2018 |
| 76766 | Crick | Biology | 72000 | 76766 | BIO-101 | 1 | Summer | 2017 |
| 76766 | Crick | Biology | 72000 | 76766 | BIO-301 | 1 | Summer | 2018 |
| 83821 | Brandt | Comp. Sci. | 92000 | 83821 | CS-190 | 1 | Spring | 2017 |
| 83821 | Brandt | Comp. Sci. | 92000 | 83821 | CS-190 | 2 | Spring | 2017 |
| 83821 | Brandt | Comp. Sci. | 92000 | 83821 | CS-319 | 2 | Spring | 2018 |
| 98345 | Kim | Elec. Eng. | 80000 | 98345 | EE-181 | 1 | Spring | 2017 |

Join Operation (Theta Join)

- The **join** operation allows us to combine a select operation and a Cartesian-Product operation into a single operation.
- Consider relations $r (R)$ and $s (S)$. Let “theta” be a predicate on attributes in the schema R “union” S . The join operation $r \bowtie_{\theta} s$ is defined as follows:

$$r \bowtie_{\theta} s = \sigma_{\theta} (r \times s)$$

- Thus

$$\sigma_{instructor.id = teaches.id} (instructor \times teaches)$$

- Can equivalently be written as

$$instructor \bowtie_{instructor.id = teaches.id} teaches.$$

Natural Join

- A special case where equality predicate holds on all attributes of same name.

| <i>Employee</i> | | |
|-----------------|-------|-----------------|
| Name | Empld | DeptName |
| Harry | 3415 | Finance |
| Sally | 2241 | Sales |
| George | 3401 | Finance |
| Harriet | 2202 | Sales |
| Mary | 1257 | Human Resources |

| <i>Dept</i> | |
|-------------|---------|
| DeptName | Manager |
| Finance | George |
| Sales | Harriet |
| Production | Charles |

| <i>Employee ⋈ Dept</i> | | | |
|------------------------|-------|----------|---------|
| Name | Empld | DeptName | Manager |
| Harry | 3415 | Finance | George |
| Sally | 2241 | Sales | Harriet |
| George | 3401 | Finance | George |
| Harriet | 2202 | Sales | Harriet |

Note that neither the employee named Mary nor the Production department appear in the result.

Rename Operation

- Allows us to refer to a relation by more than one name.
- Example:

$$\rho_x(E)$$

returns the expression E under the name X

- If a relational-algebra expression E has arity n , then


$$\rho_{x(A_1, A_2, \dots, A_n)}(E)$$

returns the result of expression E under the name X ,
and with the attributes renamed to A_1, A_2, \dots, A_n .

Find the highest salary in the university

- Steps to reach the solution:

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 12121 | Wu | Finance | 90000 |
| 15151 | Mozart | Music | 40000 |
| 22222 | Einstein | Physics | 95000 |
| 32343 | El Said | History | 60000 |
| 33456 | Gold | Physics | 87000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 58583 | Califieri | History | 62000 |
| 76543 | Singh | Finance | 80000 |
| 76766 | Crick | Biology | 72000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 98345 | Kim | Elec. Eng. | 80000 |



We find all salaries that are lesser than some existing salary.

| <i>salary</i> |
|---------------|
| 65000 |
| 90000 |
| 40000 |
| 60000 |
| 87000 |
| 75000 |
| 62000 |
| 72000 |
| 80000 |
| 92000 |



95000

The only salary left out in step2 is the max salary.

Find the highest salary in the university

- Compute instructor x instructor
- Compute

$$\Pi_{instructor.salary} (\sigma_{instructor.salary < d.salary} (instructor \times \rho_d (instructor)))$$

- Compute

$$\Pi_{salary} (instructor) - \Pi_{instructor.salary} (\sigma_{instructor.salary < d.salary} (instructor \times \rho_d (instructor)))$$

We use rename operation to distinguish the two salary attributes.

Relational Algebra

- A **basic expression** in the relational algebra consists of either one of the following:
 - A relation in the database
 - A constant relation
- Let E_1 and E_2 be **relational-algebra expressions**; the following are all relational-algebra expressions:
 - $E_1 \cup E_2$
 - $E_1 - E_2$
 - $E_1 \times E_2$
 - $\sigma_p(E_1)$, P is a predicate on attributes in E_1
 - $\Pi_s(E_1)$, S is a list consisting of some of the attributes in E_1
 - $\rho_x(E_1)$, x is the new name for the result of E_1

Problem

Consider the following relational schema.

```
Students(rollno: integer, sname: string)
Courses(courseno: integer, cname: string)
Registration(rollno: integer, courseno: integer, percent: real)
```

Is the following relational algebra expression equivalent to the English sentence given below:

"Find the distinct names of all students who score more than 90% in the course numbered 107"

$$\Pi_{\text{sname}}(\sigma_{\text{courseno}=107 \wedge \text{percent}>90} (\text{Registration} \bowtie \text{Students}))$$

Problem

Consider the relational schema given below, where *eld* of the relation *dependent* is a foreign key referring to *empId* of the relation *employee*. Assume that every employee has at least one associated dependent in the *dependent* relation.

employee (*empId*, *empName*, *empAge*)
dependent (*depId*, *eId*, *depName*, *depAge*)

Consider the following relational algebra query:

$\Pi_{\text{empId}}(\text{employee}) - \Pi_{\text{empId}}(\text{employee} \bowtie_{(\text{empId} = \text{eId}) \wedge (\text{empAge} \leq \text{depAge})} \text{dependent})$

The above query evaluates to the set of *empIds* of employees whose age is greater than that of

- (A) some dependent.
- (B) all dependents.
- (C) some of his/her dependents
- (D) all of his/her dependents.

Problem

Consider the relational schema given below, where *eld* of the relation *dependent* is a foreign key referring to *empId* of the relation *employee*. Assume that every employee has at least one associated dependent in the *dependent* relation.

employee (*empId*, *empName*, *empAge*)
dependent (*depId*, *eId*, *depName*, *depAge*)

Consider the following relational algebra query:

$\Pi_{\text{empId}}(\text{employee}) - \Pi_{\text{empId}}(\text{employee} \bowtie_{(\text{empId} = \text{eId}) \wedge (\text{empAge} \leq \text{depAge})} \text{dependent})$

The above query evaluates to the set of *empIds* of employees whose age is greater than that of

- (A) some dependent.
- (B) all dependents.
- (C) some of his/her dependents
- (D) **all of his/her dependents.**

Fundamental Operations

The select, project, rename, set difference, cartesian product, and union are sufficient to express queries!

Extended Operations

- Set Intersection

- Find the set of all courses taught in both the Fall 2017 and the Spring 2018 semesters

$$\Pi_{course_id} (\sigma_{semester="Fall" \wedge year=2017}(section)) \cap \Pi_{course_id} (\sigma_{semester="Spring" \wedge year=2018}(section))$$

Result

| |
|------------------|
| <i>course_id</i> |
| CS-101 |

- Assignment

- Find all instructors in the “Physics” and Music department.

$Physics \leftarrow \sigma_{dept_name="Physics"}(instructor)$

$Music \leftarrow \sigma_{dept_name="Music"}(instructor)$

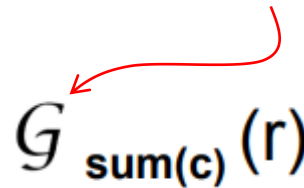
$Physics \cup Music$

Aggregate Functions

- Improves ease of use
- Most common functions:

| Functions |
|-----------|
| sum |
| max |
| min |
| avg |
| count |

Calligraphic G Notation



The diagram shows the calligraphic G notation $G_{\text{sum}(c)}(r)$. A red arrow points from the text 'Calligraphic G Notation' to the 'G' in the notation.

Compute sum of all values of attribute c on relation r.

Summary

- Relational Algebra is very much like normal algebra ($x - y$). We use relations instead of numbers as basic expressions.
- Basis for commercial query languages such as SQL.
- Three major components:
 - Fundamental Operations
 - Extended Operations
 - Aggregate Functions

Revision

| Table A | | |
|---------|--------|-----|
| Id | Name | Age |
| ----- | | |
| 12 | Arun | 60 |
| 15 | Shreya | 24 |
| 99 | Rohit | 11 |

| Table B | | |
|---------|--------|-----|
| Id | Name | Age |
| ----- | | |
| 15 | Shreya | 24 |
| 25 | Hari | 40 |
| 98 | Rohit | 20 |
| 99 | Rohit | 11 |

| Table C | | |
|---------|-------|------|
| Id | Phone | Area |
| ----- | | |
| 10 | 2200 | 02 |
| 99 | 2100 | 01 |

How many rows will this query produce?

$(A \cup B) \bowtie_{A.Id > 40 \vee C.Id < 15} C$

Revision

- 7 Rows

| Id | Name | Age | Id | Phone | Area |
|-------|--------|-----|----|-------|------|
| ----- | | | | | |
| 12 | Arun | 60 | 10 | 2200 | 02 |
| 15 | Shreya | 24 | 10 | 2200 | 02 |
| 99 | Rohit | 11 | 10 | 2200 | 02 |
| 25 | Hari | 40 | 10 | 2200 | 02 |
| 98 | Rohit | 20 | 10 | 2200 | 02 |
| 99 | Rohit | 11 | 99 | 2100 | 01 |
| 98 | Rohit | 20 | 99 | 2100 | 01 |